

# Introduction to CUDA Parallel Programming CUDA 平行計算導論

[https://ceiba.ntu.edu.tw/1092Phys8061\\_CUDA](https://ceiba.ntu.edu.tw/1092Phys8061_CUDA)

Professor Ting-Wai Chiu (趙挺偉)  
Email: [twchiu@phys.ntu.edu.tw](mailto:twchiu@phys.ntu.edu.tw)  
Physics Department  
National Taiwan University

## This lecture will cover:

- How to Perform Parallel Computation of Monte Carlo Simulation of 2D Ising Model ?
- Parallel Updating with Even-Odd (Checkerboard) Scheme
- GPU Implementation using Global Memory

# Ising Model

[https://en.wikipedia.org/wiki/Ising\\_model](https://en.wikipedia.org/wiki/Ising_model)

Consider spins on a  $D$ -dimensional lattice (torus) with energy

$$E(\sigma) = -J \sum_{\langle i, j \rangle} \sigma_i \sigma_j - B \sum_i \sigma_i$$


$$\sigma_i = \begin{cases} +1 & \text{spin up} \quad \uparrow \\ -1 & \text{spin down} \quad \downarrow \end{cases}$$

$J$  : coupling strength

$B$  : external magnetic field

where  $\sigma_i$  denotes the spin variable at the site  $i = (i_1, i_2, \dots, i_D)$ ,  $B$  is the external magnetic field,  $J$  is the coupling strength between nearest neighboring spins, and  $\langle i, j \rangle$  denotes that  $i$  and  $j$  are nearest neighboring sites.

# Ising Model on the 2D Lattice with periodic b.c.



2D Lattice with periodic b.c. = Torus

# Metropolis Algorithm for MC of Ising Model

(a) The new spin  $\sigma'_i = -\sigma_i$  is selected with prob.

distribution satisfying  $T(\sigma_i \rightarrow \sigma'_i) = T(\sigma'_i \rightarrow \sigma_i)$

(b) Compute the change in energy  $\Delta E = E(\sigma'_i) - E(\sigma_i)$

(c) If  $\Delta E \leq 0$ , then  $\sigma'_i$  is accepted, otherwise generate a random number  $r \in (0,1)$  with uniform deviate.

If  $r \leq \exp(-\Delta E / kT)$ , then  $\sigma'_i$  is accepted.

$$A(\sigma_i \rightarrow \sigma'_i) = \min \left( 1, \frac{\exp[-E(\sigma') / kT]}{\exp[-E(\sigma) / kT]} \right)$$

(d) Repeat the same procedure for another spin.

# How to Perform Parallel Updating (Simulation) with Metropolis Algorithm ?

# How to Perform Parallel Updating (Simulation) with Metropolis Algorithm ?

One way is to use Even-Odd (Checkerboard) Scheme

# How to Perform Parallel Updating (Simulation) with Metropolis Algorithm ?


One way is to use **Even-Odd** (Checkerboard) Scheme

Group the lattice sites into 2 sets, **even** sites and **odd** sites, the neighboring sites of any **even** site are **odd** sites, and vice versa. Then all **even/odd** sites can be updated in parallel. Thus the update can be first performed on all **even** sites in parallel, then go on to update all **odd** sites in parallel. This completes one sweep. (Note that even-odd scheme only works for models with the nearest neighbor coupling.)




# Even-Odd (Checkerboard) Scheme

Consider the 4x4 lattice




# Even-Odd (Checkerboard) Scheme

Consider the 4x4 lattice




# Even-Odd (Checkerboard) Scheme

Consider the 4x4 lattice




# Even-Odd (Checkerboard) Scheme

Consider the 4x4 lattice



# Even-Odd (Checkerboard) Scheme

Consider the 4x4 lattice



12 13 14 15  
08 09 10 11  
04 05 06 07  
00 01 02 03



Even sites

08 10 13 15  
00 02 05 07

**From even space to full space**

$x = (2*iE) \% Nx$ ;  $y = ((2*iE)/Nx) \% Nx$ ;


$parity = (x+y) \% 2$ ;

$x = x + parity$ ;

$i = x + Nx*y$ ;

# Even-Odd (Checkerboard) Scheme

Consider the 4x4 lattice



12 13 14 15  
08 09 10 11  
04 05 06 07  
00 01 02 03



Odd sites

09 11 12 14  
01 03 04 06



**From odd space to full space**

$x = (2*i0)\%Nx;$   $y = ((2*i0)/Nx)\%Nx;$


$parity = (x+y+1)\%2;$

$x = x+parity;$

$i = x+Nx*y;$

# Even-Odd (Checkerboard) Scheme

Consider the 4x4 lattice



12 13 14 15  
08 09 10 11  
04 05 06 07  
00 01 02 03



Odd sites



Even sites

09 11 12 14  
01 03 04 06  
08 10 13 15  
00 02 05 07

**From even space to full space**

```
x = (2*iE)%Nx; y = ((2*iE)/Nx)%Nx;
parity = (x+y)%2;
x = x+parity;
i = x+Nx*y;
```

**From odd space to full space**

```
x = (2*iO)%Nx; y = ((2*iO)/Nx)%Nx;
parity = (x+y+1)%2;
x = x+parity;
i = x+Nx*y;
```

# Even-Odd (Checkerboard) Scheme

```
#include <stdio.h>

void main()
{
    int Nx=4;
    int x, y, i e, i , parity;
    int sites=Nx*Nx;
    int half_sites = sites/2;

    for(i e=0; i e<half_sites; i e++) {
        x=(2*i e)%Nx;
        y=((2*i e)/Nx)%Nx;
        parity=(x+y)%2;
        x = x + parity;
        i = x + y*Nx;
        printf("%2d %2d %2d %2d\n", i e, x, y, i );
    }
    for(i o=0; i o<half_sites; i o++) {
        x=(2*i o)%Nx;
        y=((2*i o)/Nx)%Nx;
        parity=(x+y+1)%2;
        x = x + parity;
        i = x + y*Nx;
        printf("%2d %2d %2d %2d\n", i o, x, y, i );
    }
}
```



# Even-Odd (Checkerboard) Scheme

```
#include <stdio.h>
```

```
void main()
```

```
{
```

```
    int Nx=4;
```

```
    int x, y, i e, i , parity;
```

```
    int sites=Nx*Nx;
```

```
    int half_sites = sites/2;
```

```
    for(i e=0; i e<half_sites; i e++) {
```

```
        x=(2*i e)%Nx;
```

```
        y=((2*i e)/Nx)%Nx;
```

```
        parity=(x+y)%2;
```

```
        x = x + parity;
```

```
        i = x + y*Nx;
```

```
        printf("%2d %2d %2d %2d\n", i e, x, y, i );
```

```
    }
```

```
    for(i o=0; i o<half_sites; i o++) {
```

```
        x=(2*i o)%Nx;
```

```
        y=((2*i o)/Nx)%Nx;
```

```
        parity=(x+y+1)%2;
```


```
        x = x + parity;
```

```
        i = x + y*Nx;
```

```
        printf("%2d %2d %2d %2d\n", i o, x, y, i );
```

```
    }
```

```
} // The complete code at twqcd80: /home/cuda_lecture_2021/Ising2D/index_eo.c
```



i e	x	y	i
=====			
0	0	0	0
1	2	0	2
2	1	1	5
3	3	1	7
4	0	2	8
5	2	2	10
6	1	3	13
7	3	3	15

i o	x	y	i
=====			
0	1	0	1
1	3	0	3
2	0	1	4
3	2	1	6
4	1	2	9
5	3	2	11
6	0	3	12
7	2	3	14

# Parallel Updating (Simulation) with Even-Odd (Checkerboard) Scheme


```
for(int ie=0; ie<ns/2; ie++) {           // update even sites
    x = (2*ie)%nx;
    y = ((2*ie)/nx)%nx;
    parity=(x+y)%2;
    x = x + parity;
    i = x + y*nx;
    old_spin = spin[i];
    new_spin = -old_spin;
    k1 = fw[x] + y*nx;           // right
    k2 = x + fw[y]*nx;           // top
    k3 = bw[x] + y*nx;           // left
    k4 = x + bw[y]*nx;           // bottom
    spins = spin[k1] + spin[k2] + spin[k3] + spin[k4];
    de = -(new_spin - old_spin)*(spins + B);
    if((de <= 0.0) || (gsl_rng_uniform(rng) < exp(-de/T))) {
// if (new energy <= old energy) or (r < exp(-dE/kT)) accept the new spin
        spin[i] = new_spin;       // accept the new spin;
    }
}
```

# Parallel Updating (Simulation) with Even-Odd (Checkerboard) Scheme

```
for(int io=0; io<ns/2; ie++) {           // update odd sites
    x = (2*io)%nx;
    y = ((2*io)/nx)%nx;
    parity=(x+y+1)%2;
    x = x + parity;
    i = x + y*nx;
    old_spin = spin[i];
    new_spin = -old_spin;
    k1 = fw[x] + y*nx;    // right
    k2 = x + fw[y]*nx;    // top
    k3 = bw[x] + y*nx;    // left
    k4 = x + bw[y]*nx;    // bottom
    spins = spin[k1] + spin[k2] + spin[k3] + spin[k4];
    de = -(new_spin - old_spin)*(spins + B);
    if((de <= 0.0) || (gsl_rng_uniform(rng) < exp(-de/T))) {
// if (new energy <= old energy) or (r < exp(-dE/kT)) accept the new spin
        spin[i] = new_spin;           // accept the new spin;
    }
} //Complete code at twqcd80: /home/cuda_lecture_2021/Ising2D/ising2d_cpu_eo.c
```

# GPU Implementation with Global Memory

Lattice size: 16 x 16  
gridDim = (4, 4)  
blockDim = (2, 4)



# GPU Implementation with Global Memory

Some considerations:

- GPU kernel can perform several sweeps at each call
- To use CPU or GPU to generate RNG ?
- To put the entire simulation in the GPU ?
- Instead of computing  $\exp(-\Delta E/kT)$ , its value can be looked up from a table which is prepared before the start of the simulation, since the number of possible values of  $\exp(-\Delta E/kT)$  are limited, due to the discreteness of the spin.