

Introduction to CUDA Parallel Programming CUDA 平行計算導論

https://ceiba.ntu.edu.tw/1092Phys8061_CUDA

Professor Ting-Wai Chiu (趙挺偉)
Email: twchiu@phys.ntu.edu.tw
Physics Department
National Taiwan University

This lecture will cover:

- CUDA Libraries: cuBLAS, cuFFT,...
- cuBLAS – Saxpy: $|B\rangle \leftarrow \alpha |A\rangle + |B\rangle$
- cuBLAS – Sgemm: $C \leftarrow \alpha A \cdot B + \beta C$

The cuBLAS library is an implementation of BLAS (Basic Linear Algebra Subprograms) on top of the NVIDIA CUDA runtime. It allows the user to access the computational resources of NVIDIA GPU.

<https://docs.nvidia.com/cuda/cublas/index.html>

BLAS – Basic Linear Algebra Subroutines

<http://www.netlib.org/blas/>

BLAS Routines

LEVEL 1 SINGLE

- [SROTG](#) - setup Givens rotation
- [SROTMG](#) - setup modified Givens rotation
- [SROT](#) - apply Givens rotation
- [SROTM](#) - apply modified Givens rotation
- [SSWAP](#) - swap x and y
- [SSCAL](#) - $x = a * x$
- [SCOPY](#) - copy x into y
- [SAXPY](#) - $y = a * x + y$
- [SDOT](#) - dot product
- [SDSDOT](#) - dot product with extended precision accumulation
- [SNRM2](#) - Euclidean norm
- [SCNRM2](#) - Euclidean norm
- [SASUM](#) - sum of absolute values
- [ISAMAX](#) - index of max abs value

BLAS – Basic Linear Algebra Subroutines

<http://www.netlib.org/blas/>

BLAS Routines

LEVEL 1 DOUBLE

[DROTG](#) - setup Givens rotation

[DROTMG](#) - setup modified Givens rotation

[DROT](#) - apply Givens rotation

[DROTM](#) - apply modified Givens rotation

[DSWAP](#) - swap x and y

[DSCAL](#) - $x = a * x$

[DCOPY](#) - copy x into y

[DAXPY](#) - $y = a * x + y$

[DDOT](#) - dot product

[DSDOT](#) - dot product with extended precision accumulation

[DNRM2](#) - Euclidean norm

[DZNRM2](#) - Euclidean norm

[DASUM](#) - sum of absolute values

[IDAMAX](#) - index of max abs value

BLAS – Basic Linear Algebra Subroutines

<http://www.netlib.org/blas/>

BLAS Routines

LEVEL 1 COMPLEX

[CROTG](#) - setup Givens rotation

[CSROT](#) - apply Givens rotation

[CSWAP](#) - swap x and y

[CSCAL](#) - $x = a * x$

[CSSCAL](#) - $x = a * x$

[CCOPY](#) - copy x into y

[CAXPY](#) - $y = a * x + y$

[CDOTU](#) - dot product

[CDOTC](#) - dot product, conjugating the first vector

[SCASUM](#) - sum of absolute values

[ICAMAX](#) - index of max abs value

BLAS – Basic Linear Algebra Subroutines

<http://www.netlib.org/blas/>

BLAS Routines

LEVEL 1 DOUBLE COMPLEX

[ZROTG](#) - setup Givens rotation

[ZDROTF](#) - apply Givens rotation

[ZSWAP](#) - swap x and y

[ZSCAL](#) - $x = a * x$

[ZDSCAL](#) - $x = a * x$

[ZCOPY](#) - copy x into y

[ZAXPY](#) - $y = a * x + y$

[ZDOTU](#) - dot product

[ZDOTC](#) - dot product, conjugating the first vector

[DZASUM](#) - sum of absolute values

[IZAMAX](#) - index of max abs value

BLAS – Basic Linear Algebra Subroutines

<http://www.netlib.org/blas/>

BLAS Routines

LEVEL 2 S(Single) / D(Double) / C(Complex) / Z(Double Complex)

SGEMV - **matrix vector** multiply

SGBMV - banded matrix vector multiply

SSYMV - symmetric matrix vector multiply

SSBMV - symmetric banded matrix vector multiply

SSPMV - symmetric packed matrix vector multiply

STRMV - triangular matrix vector multiply

STBMV - triangular banded matrix vector multiply

STPMV - triangular packed matrix vector multiply

STRSV - solving triangular matrix problems

STBSV - solving triangular banded matrix problems

STPSV - solving triangular packed matrix problems

SGER - performs the rank 1 operation $A := \alpha * x * y' + A$

SSYR - performs the symmetric rank 1 operation $A := \alpha * x * x' + A$

SSPR - symmetric packed rank 1 operation $A := \alpha * x * x' + A$

SSYR2 - performs the symmetric rank 2 operation, $A := \alpha * x * y' + \alpha * y * x' + A$

SSPR2 - performs the symmetric packed rank 2 operation, $A := \alpha * x * y' + \alpha * y * x' + A$

BLAS – Basic Linear Algebra Subroutines

<http://www.netlib.org/blas/>

BLAS Routines

Level 3 **Z(Double Complex) / C(Complex) / D(Double) / S(Single)**

ZGEMM - **matrix matrix** multiply

ZSYMM - symmetric matrix matrix multiply

ZHEMM - hermitian matrix matrix multiply

ZSYRK - symmetric rank-k update to a matrix

ZHERK - hermitian rank-k update to a matrix

ZSYR2K - symmetric rank-2k update to a matrix

ZHER2K - hermitian rank-2k update to a matrix

ZTRMM - triangular matrix matrix multiply

ZTRSM - solving triangular matrix with multiple right hand sides

cuBLAS

The cuBLAS library is an implementation of BLAS (Basic Linear Algebra Subprograms) on top of the NVIDIA CUDA runtime. It allows the user to access the computational resources of NVIDIA GPU.

<https://docs.nvidia.com/cuda/cublas/index.html>

BLAS Sgemm $C \leftarrow \alpha A \cdot B + \beta C$

cuBLAS **cublas**Sgemm for single GPU

cuBLAS **cublasXt**Sgemm for multi-GPUs
(Note that only some routines are available for multi-GPUs)

cuBLAS API

The cuBLAS Library exposes three sets of API:

- The [cuBLAS API](#), which is simply called cuBLAS API (starting with CUDA 6.0)
- The [CUBLASXT API](#) (starting with CUDA 6.0)
- The [cuBLASLt API](#) (starting with CUDA 10.1)

To use the cuBLAS API, **the application must allocate the required matrices and vectors in the GPU memory space**, fill them with data, call the sequence of desired cuBLAS functions, and then upload the results from the GPU memory space back to the host. The cuBLAS API also provides helper functions for writing and retrieving data from the GPU.

To use the CUBLASXT API, **the application must keep the data on the Host and the Library will take care of dispatching the operation to one or multiple GPUs present in the system, depending on the user request.**

If you can find your needed routine in CUBLASXT API, then use it rather than the corresponding one in cuBLAS API.

The cuBLASLt API is a lightweight library dedicated to General Matrix-to-matrix Multiply (GEMM) operations with a new flexible API.

SAXPY($y \leftarrow a * x + y$), simple kernel with 1-GPU

```
__global__ void Saxpy(const float alpha, const float* A, float* B, long N)
{
    long i = blockDim.x * blockIdx.x + threadIdx.x;

    while (i < N) {
        B[i] += alpha*A[i];
        i += blockDim.x * gridDim.x; // go to the next grid
    }
    __syncthreads();
}
```

// complete code in [twqcd80:/home/cuda_lecture_2021/Saxpy_1GPU/Saxpy_1gpu.cu](#)

SAXPY ($y \leftarrow a*x + y$), cuBLAS with 1-GPU


```
#include "cublas_v2.h"           // header for CUBLAS

int main(void){
    ...
    cublasHandle_t handle;        // CUBLAS context
    ...
    cublasCreate(&handle);        // initialize CUBLAS context
    cublasSetVector(N, sizeof(float), h_A, 1, d_A, 1); // copy h_A to d_A
    cublasSetVector(N, sizeof(float), h_B, 1, d_B, 1);

    cublasSaxpy(handle, N, &alpha, d_A, 1, d_B, 1);    // B <- alpha*A + B

    cublasGetVector(N, sizeof(float), d_B, 1, h_C, 1); // copy d_B to h_C
    ...
    cublasDestroy(handle);        // destroy cublas context
}
// twqcd80: /home/cuda_lecture_2021/Saxpy_1GPU/Saxpy_1gpu_cublas.cu
```

Unified Memory in CUDA

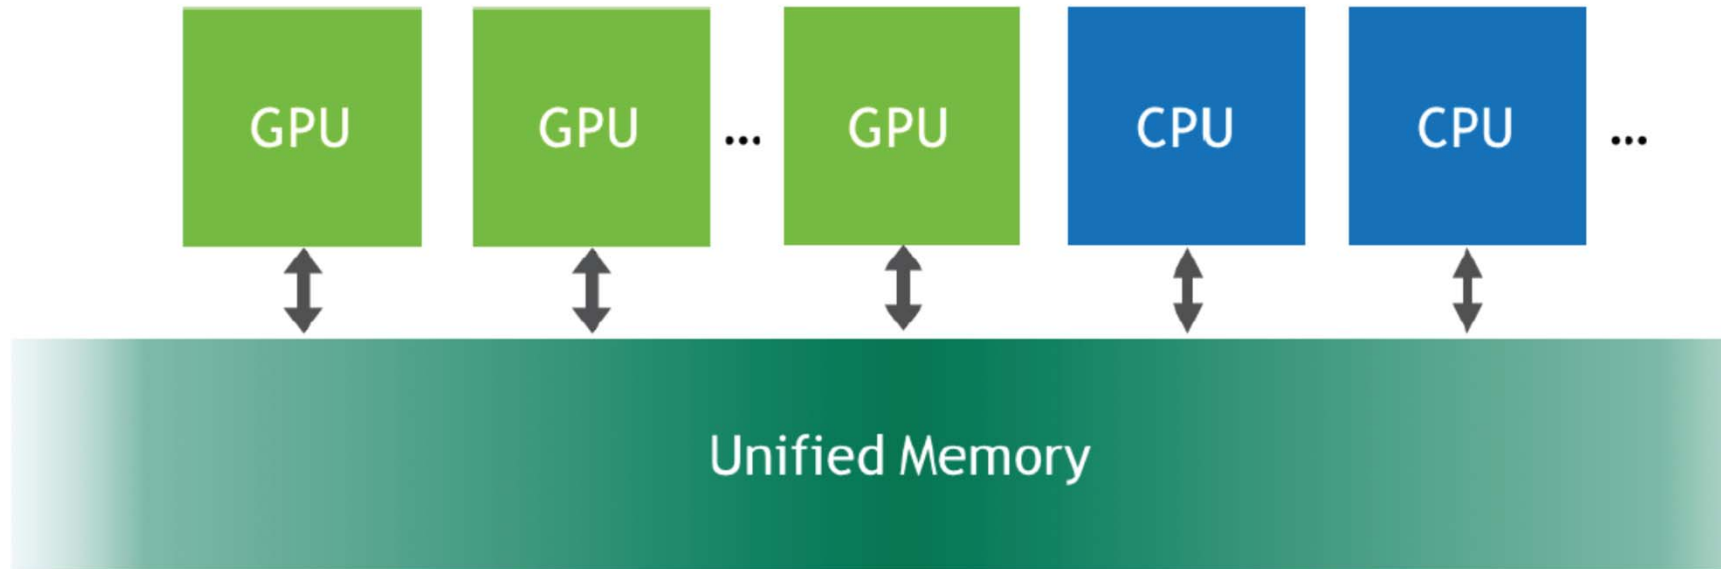


Unified Memory in CUDA is a single address space accessible from any processor in a system. This hardware/software technology allows applications to allocate data that can be read or written from code running on either CPUs or GPUs.

~~malloc()
cudaMalloc()~~

cudaMallocManaged()

Unified Memory in CUDA



Unified Memory in CUDA is a single address space accessible from any processor in a system. This hardware/software technology allows applications to allocate data that can be read or written from code running on either CPUs or GPUs.

~~cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice)
cudaMemcpy(h_A, d_A, size, cudaMemcpyDeviceToHost)~~

SAXPY, cuBLAS with 1-GPU, unified memory

```
#include "cublas_v2.h"           // header for CUBLAS

int main(void){
    ...
    cublasHandle_t handle;        // CUBLAS context
    ...
    cudaMallocManaged(&h_A, size); // Allocate h_A and h_B in unified memory
    cudaMallocManaged(&h_B, size);
    ...
    cublasCreate(&handle);         // initialize CUBLAS context
    cublasSaxpy(handle, N, &alpha, h_A, 1, h_B, 1); // B <- alpha*A + B
    ...
    cublasDestroy(handle);        // destroy cublas context
}

// twqcd80: /home/cuda_lecture_2021/Saxpy_1GPU/Saxpy_1gpu_cublas_umem.cu
```

Sgemm [$C \leftarrow \alpha A \cdot B + \beta C$], multi-GPUs, unified memory

```
__global__ void sgemm_umem(int n, float alpha, const float *A, const float *B,
                           const float beta, float *C, const int NGPU, const int cpu_thread_id)
{
    int offset = n/NGPU*cpu_thread_id;
    int i = blockDim.x * blockIdx.x + threadIdx.x + offset;
    for(int j = 0; j < n; ++j) {
        float prod = 0.0f;
        for(int k = 0; k < n; ++k) prod += A[i*n+k]*B[k*n+j];    // A(i,k)*B(k,j)
        C[i*n + j] = alpha*prod + beta*C[i*n + j];
    }
    __syncthreads();
}

int main(void)
{
    ...
    cudaMallocManaged((void**)&h_A, size);    // allocate unified memory
    cudaMallocManaged((void**)&h_B, size);
    cudaMallocManaged((void**)&h_C, size);
    ...
    #pragma omp parallel private(cpu_thread_id){
        cpu_thread_id = omp_get_thread_num();
        cudaSetDevice(Dev[cpu_thread_id]);
        sgemm_umem<<<blocks, threads>>>(N, alpha, h_A, h_B, beta, h_C, NGPU, cpu_thread_id);
        cudaDeviceSynchronize();
    }
    ...
} // The complete code at twqcd80: /home/cuda_lecture_2021/Sgemm_NGPU/Sgemm_umem.cu
```


Sgemm, cuBLAS with multi-GPUs, unified memory

```
#include "cublasXt.h"           // header for cublasXT

int main(void){
    ...
    cublasXtHandle_t handle;     // cublasXt context
    ...
    cudaMallocManaged(&h_A, size); // Allocate h_A, h_B, h_C in unified memory
    cudaMallocManaged(&h_B, size);
    cudaMallocManaged(&h_C, size);
    ...
    cublasXtCreate(&handle);      // initialize cublasXt context
    cublasXtDeviceSelect(handle, NGPU, Dev); // Select devices for CUBLASXT
    cublasXtSgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N, N, N, N, &alpha, h_A, N, h_B, N,
                  &beta, h_C, N); // C <- alpha*A.B + beta*C
    ...
    cublasXtDestroy(handle);      // destroy cublasXt context
}

// twqcd80: /home/cuda_lecture_2021/Sgemm_NGPU/Sgemm_ngpu_cublas.cu
```