Introduction to CUDA Parallel Programming
https://ceiba.ntu.edu.tw/1092Phys8061_CUDA
Instructions for using the GPU nodes for this course.

Three GPU nodes which have been assigned to this course. Each node contains 2 GPUs, Nvidia GTX1060, with CUDA 10.2 installed in /usr/local/nvidia

You can login each of these 3 nodes via ssh,

```
ssh -p 50080 username@twcp1.phys.ntu.edu.tw     (twqcd80)
ssh -p 50074 username@twcp1.phys.ntu.edu.tw     (twqcd74)
ssh -p 50218 username@twcp1.phys.ntu.edu.tw     (twqcd218)
```

If your student ID starts with a letter, e.g., r01234567, then your user name is the same as your student ID. On the other hand, if your student ID ends with the letter s, e.g., 40123456s, then your user name is s40123456. The initial password is the same as the user name. You should change your password immediately at the first login.

To ssh another node from any node, the port number is not needed,
for example, from twqcd80 ssh to twqcd74

username@twqcd80:~$ ssh twqcd74

Before you run your GPU job at any node, you should use **nvidia-smi** to check whether anyone is using the GPUs, and to select the free GPUs. In case if both GPUs are being used, you can ssh to another node. Since **these 3 nodes do not have a common file system**, you need to transfer your files from one node to another node, e.g., using **scp**

In twqcd80:**/home/cuda_lecture_2021** contains **the sample codes for this course**.